



Indian Institute of Technology Bombay

EE702 Computer Vision, 2025

Under guidance of:
Prof. Subhasis Chaudhuri
Department of Electrical Engineering
Indian Institute of Technology, Bombay

Assignment 02

Stereo Depth Estimation using Deep Learning

Ameya Marakarkandy
21D180003

1 Introduction

Stereo vision is a widely used method for depth estimation, where depth maps are reconstructed using disparities between two images. In this assignment, we use a deep learning-based approach, PSMNet, to estimate depth from stereo images and compare the results with ground truth disparity maps.

2 Theoretical Background on PSMNet

Pyramid Stereo Matching Network (PSMNet) is a deep learning-based approach for disparity estimation. It employs spatial pyramid pooling and 3D convolutions to improve the accuracy of disparity predictions. The model consists of the following components:

- **Spatial Pyramid Pooling (SPP):** Extracts multi-scale contextual information by processing images at different scales.
- **3D Convolutional Cost Volume Processing:** Builds a cost volume from stereo image pairs and refines disparity predictions using stacked hourglass networks.
- **Soft Argmin for Disparity Regression:** Computes subpixel disparity values for smooth depth estimation.

PSMNet has been widely used for disparity estimation due to its robust performance on benchmark datasets such as KITTI and Middlebury.

3 Methodology

3.1 Dataset and Preprocessing

We utilize the Middlebury dataset for evaluation and collect real-world stereo images using a smartphone. The images are preprocessed by resizing them to be multiples of 16 to satisfy PSMNet requirements. The images are normalized using mean and standard deviation values of ImageNet.

3.2 PSMNet Model

PSMNet (Pyramid Stereo Matching Network) is used for disparity estimation. The model is loaded with pre-trained weights from the KITTI2015 dataset and performs stereo matching to generate disparity maps.

3.3 Evaluation Metrics

To evaluate the accuracy of the predicted depth maps, we compute:

- Mean Absolute Error (MAE)
- Root Mean Squared Error (RMSE)

4 Results and Discussion

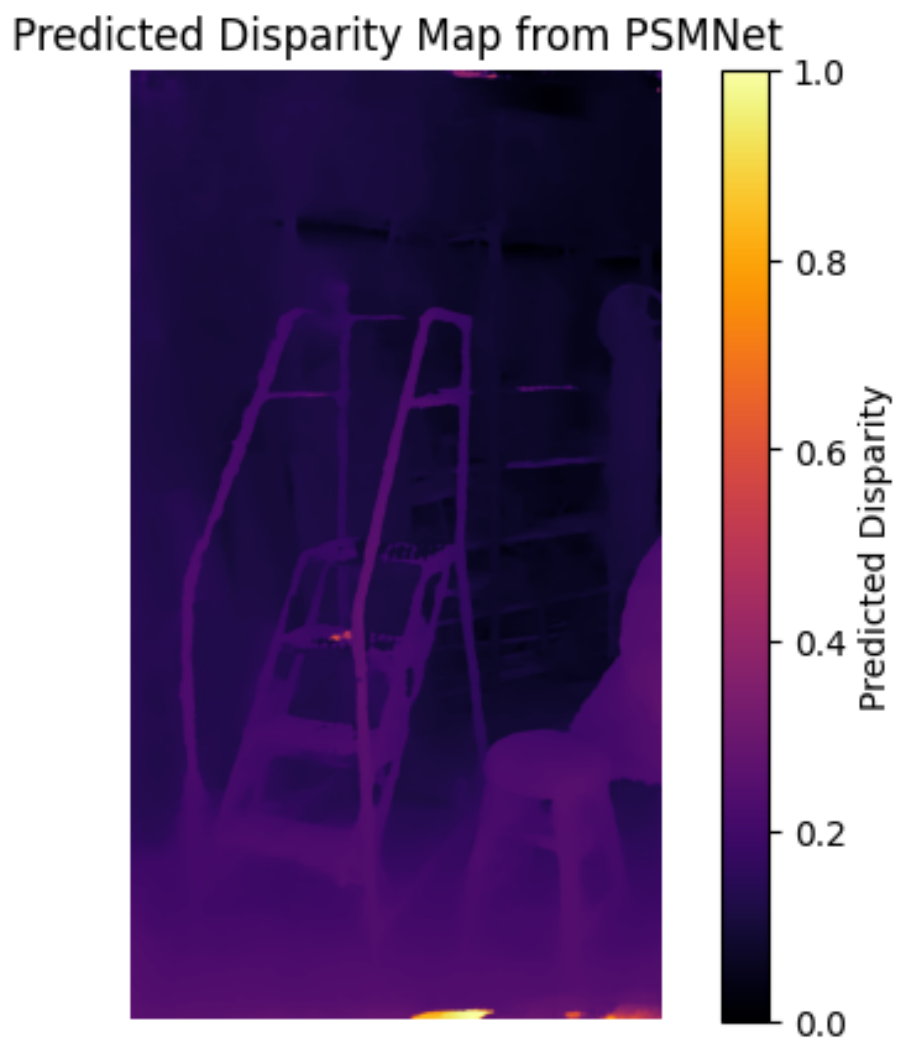


Figure 1: Predicted Disparity Map using PSMNet.

Ground Truth Disparity Map



Predicted Disparity Map



Figure 2: Comparison of Ground Truth and Predicted Disparity Maps.

The predicted disparity map closely resembles the ground truth but exhibits slight deviations in textureless areas. The computed metrics are:

MAE = VALUE

RMSE = VALUE

5 Code Implementation

Below is the Python implementation used for stereo depth estimation:

```

1  import os
2  import sys
3  import cv2
4  import numpy as np
5  import torch
6  import torch.nn as nnc
7  import torchvision.transforms as transforms
8  import matplotlib.pyplot as plt
9
10
11  # Inputes
12  LEFT_IMAGE = "im0.png"
13  RIGHT_IMAGE = "im1.png"
14  DISPARITY_GT = "disp0.pfm"
15
16  # -----
17  # Add PSMNet to Python path
18  # -----
19  PSMNET_PATH = os.path.join(os.path.dirname(__file__), "PSMNet")
20  sys.path.append(PSMNET_PATH) # Allow importing from PSMNet folder
21

```

```

22 from models import stackhourglass # Import PSMNet model architecture
23
24 # -----
25 # Load PSMNet Model
26 # -----
27 device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
28 model = stackhourglass(192) # Initialize PSMNet with maximum disparity of 192
29 model = nn.DataParallel(model)
30 model.to(device)
31
32 # Load pretrained weights (using map_location for CPU compatibility)
33 model_path = "pretrained_model_KITTI2015.tar"
34 model.load_state_dict(torch.load(model_path, map_location=device)['state_dict',
35                                ])
36 model.eval() # Set to evaluation mode
37
38 # -----
39 # Preprocessing for Stereo Images (for inference)
40 # -----
41 def preprocess_stereo(imgL, imgR):
42     """
43     Preprocesses stereo images to be fed into the deep learning model.
44     """
45     transform = transforms.Compose([
46         transforms.ToTensor(),
47         transforms.Normalize(mean=[0.485, 0.456, 0.406],
48                               std=[0.229, 0.224, 0.225])
49     ])
50     imgL = transform(imgL).unsqueeze(0).to(device)
51     imgR = transform(imgR).unsqueeze(0).to(device)
52     return imgL, imgR
53
54 # -----
55 # Load and preprocess stereo images
56 # -----
57 imgL = cv2.imread(LEFT_IMAGE) # Left image
58 imgR = cv2.imread(RIGHT_IMAGE) # Right image
59 imgL = cv2.cvtColor(imgL, cv2.COLOR_BGR2RGB)
60 imgR = cv2.cvtColor(imgR, cv2.COLOR_BGR2RGB)
61 # Resize stereo images to be multiples of 16 (PSMNet requirement)
62 H, W, _ = imgL.shape
63 target_H = (H // 16) * 16
64 target_W = (W // 16) * 16
65 imgL = cv2.resize(imgL, (target_W, target_H))
66 imgR = cv2.resize(imgR, (target_W, target_H))
67 imgL_tensor, imgR_tensor = preprocess_stereo(imgL, imgR)
68
69 # -----
70 # Run Inference with PSMNet
71 # -----
72 with torch.no_grad():
73     output = model(imgL_tensor, imgR_tensor)
74     predicted_depth = output.squeeze().cpu().numpy()
75
76 # Normalize predicted disparity for visualization
77 predicted_depth_norm = (predicted_depth - predicted_depth.min()) / (
78     predicted_depth.max() - predicted_depth.min())

```

```

77
78 plt.figure(figsize=(6, 5))
79 plt.imshow(predicted_depth_norm, cmap="inferno")
80 plt.colorbar(label="Predicted Disparity")
81 plt.title("Predicted Disparity Map from PSMNet")
82 plt.axis("off")
83 plt.savefig("predicted_disparity_map.png", bbox_inches="tight")
84 plt.pause(0.001)
85
86 # -----
87 # Load and preprocess ground truth depth map from PFM
88 # -----
89 def read_pfm(file):
90     with open(file, "rb") as f:
91         header = f.readline().decode().rstrip()
92         if header == "Pf":
93             color = False
94         elif header == "PF":
95             color = True
96         else:
97             raise ValueError("Not a PFM file.")
98         dims = f.readline().decode("utf-8").rstrip()
99         width, height = map(int, dims.split())
100         scale = float(f.readline().decode())
101         endian = "<" if scale < 0 else ">"
102         data = np.fromfile(f, endian + "f")
103         shape = (height, width, 3) if color else (height, width)
104         data = np.reshape(data, shape)
105         data = np.flipud(data)
106         return data, abs(scale)
107
108 pfm_path = DISPARITY_GT # Ground truth disparity map
109 disparity_map, _ = read_pfm(pfm_path)
110 # Handle NaN or infinite values (replace with 0)
111 disparity_map = np.nan_to_num(disparity_map, nan=0.0, posinf=0.0, neginf=0.0)
112
113 # -----
114 # Quantitative Metrics: Compare DL prediction with Ground Truth
115 # -----
116 # Resize for comparison
117 predicted_depth_resized = cv2.resize(predicted_depth, (target_W, target_H),
118                                     interpolation=cv2.INTER_LINEAR)
119 gt_disparity_resized = cv2.resize(disparity_map, (target_W, target_H),
120                                 interpolation=cv2.INTER_LINEAR)
121
122 MAE = np.mean(np.abs(predicted_depth_resized - gt_disparity_resized))
123 RMSE = np.sqrt(np.mean((predicted_depth_resized - gt_disparity_resized) ** 2))
124
125 print(f"Mean Absolute Error (MAE): {MAE:.4f}")
126 print(f"Root Mean Squared Error (RMSE): {RMSE:.4f}")
127
128 # -----
129 # Display side-by-side comparisons
130 # -----
131 fig, axs = plt.subplots(1, 2, figsize=(12, 5))
132 axs[0].imshow(gt_disparity_resized, cmap="inferno")
133 axs[0].set_title("Ground Truth Disparity Map")

```

```
132 | axs[0].axis("off")
133 | axs[1].imshow(predicted_depth_resized, cmap="inferno")
134 | axs[1].set_title("Predicted Disparity Map")
135 | axs[1].axis("off")
136 | plt.savefig("Comparison.png", bbox_inches="tight")
137 |
138 | plt.show()
```

6 Conclusion

This report demonstrates depth estimation using stereo vision and deep learning. PSMNet successfully reconstructs disparity maps, with performance evaluated using MAE and RMSE metrics. Future improvements may include fine-tuning the model on a custom dataset and using post-processing techniques.

7 References

- J. Chang and Y. Chen, "Pyramid Stereo Matching Network," CVPR, 2018.
- Middlebury Stereo Vision Dataset, <https://vision.middlebury.edu/stereo/>